

# 3CNN, MaxPool, ReLU, Batch, SGD

March 25, 2024

```
[1]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import torch
import torchvision
import torchvision.transforms as transforms
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
```

## Prepare for Dataset

```
[2]: transform = transforms.Compose(
    [#transforms.RandomHorizontalFlip(),
     #transforms.RandomRotation(10),
     transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, #Increase ↴
                                         shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                         shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```


Files already downloaded and verified

Files already downloaded and verified

```
[3]: def imshow(img):
    img = img / 2 + 0.5
    npimg = img.numpy()
```

```
plt.imshow(np.transpose(npimg, (1, 2, 0)))
plt.show()

dataiter = iter(trainloader)
images, labels = next(dataiter)
imshow(torchvision.utils.make_grid(images))
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



cat bird horse frog

### Choose a Device

```
[4]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda:0

### Network Definition

```
[5]: class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()

        self.conv1 = nn.Conv2d(3, 32, 3, 1, 1)
        self.batch1 = nn.BatchNorm2d(32)
        self.activation1 = nn.ReLU()
        self.mpool1 = nn.MaxPool2d(kernel_size = (2,2))

        self.conv2 = nn.Conv2d(32, 16, 3, 1, 1)
        self.batch3 = nn.BatchNorm2d(16)
        self.activation2 = nn.ReLU()
        self.mpool2 = nn.MaxPool2d(kernel_size = (2,2))

        self.conv3 = nn.Conv2d(16, 64, 3, 1, 1)
        self.batch3 = nn.BatchNorm2d(64)
        self.activation3 = nn.ReLU()
        self.mpool3 = nn.MaxPool2d(kernel_size = (2,2))
```

```

        self.flat = nn.Flatten()

        self.fc1 = nn.Linear(1024,100)
        self.activation3 = nn.ReLU()
        self.fc2 = nn.Linear(100,10)

    def forward(self, x):

        x = self.activation1(self.conv1(x))
        x = self.mpool1(x)
        x = self.activation2(self.conv2(x))
        x = self.mpool2(x)
        x = self.activation3(self.conv3(x))
        x = self.mpool3(x)
        x = self.flat(x)
        x = self.fc1(self.activation3(x))
        x = self.fc2(x)

    return x

net = Net()
net.to(device)

```

```

[5]: Net(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (batch1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (activation1): ReLU()
  (mpool1): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceil_mode=False)
  (conv2): Conv2d(32, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (batch3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (activation2): ReLU()
  (mpool2): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceil_mode=False)
  (conv3): Conv2d(16, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (activation3): ReLU()
  (mpool3): MaxPool2d(kernel_size=(2, 2), stride=(2, 2), padding=0, dilation=1,
ceil_mode=False)
  (flat): Flatten(start_dim=1, end_dim=-1)
  (fc1): Linear(in_features=1024, out_features=100, bias=True)
  (fc2): Linear(in_features=100, out_features=10, bias=True)
)

```

## Optimizer and Loss Function

```
[6]: loss_func = nn.CrossEntropyLoss()

opt = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

## Training Procedure

```
[7]: avg_losses = []
epochs = 15
print_freq = 1500

for epoch in range(epochs):
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):

        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)
        opt.zero_grad()
        outputs = net(inputs)
        loss = loss_func(outputs, labels)
        loss.backward()
        opt.step()

        running_loss += loss.item()
        if i % print_freq == print_freq - 1:
            avg_loss = running_loss / print_freq
            print('[epoch: {}], i: {:5d}] avg mini-batch loss: {:.3f}'.format(
                epoch, i, avg_loss))
            avg_losses.append(avg_loss)
            running_loss = 0.0

print('Finished Training.')
```

```
[epoch: 0, i: 1499] avg mini-batch loss: 2.126
[epoch: 0, i: 2999] avg mini-batch loss: 1.729
[epoch: 0, i: 4499] avg mini-batch loss: 1.563
[epoch: 0, i: 5999] avg mini-batch loss: 1.479
[epoch: 0, i: 7499] avg mini-batch loss: 1.379
[epoch: 0, i: 8999] avg mini-batch loss: 1.345
[epoch: 0, i: 10499] avg mini-batch loss: 1.281
[epoch: 0, i: 11999] avg mini-batch loss: 1.235
[epoch: 1, i: 1499] avg mini-batch loss: 1.171
[epoch: 1, i: 2999] avg mini-batch loss: 1.135
[epoch: 1, i: 4499] avg mini-batch loss: 1.087
[epoch: 1, i: 5999] avg mini-batch loss: 1.081
[epoch: 1, i: 7499] avg mini-batch loss: 1.053
[epoch: 1, i: 8999] avg mini-batch loss: 1.050
[epoch: 1, i: 10499] avg mini-batch loss: 1.049
```


[epoch: 1, i: 11999] avg mini-batch loss: 1.003  
[epoch: 2, i: 1499] avg mini-batch loss: 0.959  
[epoch: 2, i: 2999] avg mini-batch loss: 0.932  
[epoch: 2, i: 4499] avg mini-batch loss: 0.951  
[epoch: 2, i: 5999] avg mini-batch loss: 0.944  
[epoch: 2, i: 7499] avg mini-batch loss: 0.916  
[epoch: 2, i: 8999] avg mini-batch loss: 0.933  
[epoch: 2, i: 10499] avg mini-batch loss: 0.914  
[epoch: 2, i: 11999] avg mini-batch loss: 0.923  
[epoch: 3, i: 1499] avg mini-batch loss: 0.828  
[epoch: 3, i: 2999] avg mini-batch loss: 0.847  
[epoch: 3, i: 4499] avg mini-batch loss: 0.845  
[epoch: 3, i: 5999] avg mini-batch loss: 0.874  
[epoch: 3, i: 7499] avg mini-batch loss: 0.863  
[epoch: 3, i: 8999] avg mini-batch loss: 0.882  
[epoch: 3, i: 10499] avg mini-batch loss: 0.858  
[epoch: 3, i: 11999] avg mini-batch loss: 0.869  
[epoch: 4, i: 1499] avg mini-batch loss: 0.778  
[epoch: 4, i: 2999] avg mini-batch loss: 0.800  
[epoch: 4, i: 4499] avg mini-batch loss: 0.789  
[epoch: 4, i: 5999] avg mini-batch loss: 0.824  
[epoch: 4, i: 7499] avg mini-batch loss: 0.814  
[epoch: 4, i: 8999] avg mini-batch loss: 0.789  
[epoch: 4, i: 10499] avg mini-batch loss: 0.804  
[epoch: 4, i: 11999] avg mini-batch loss: 0.817  
[epoch: 5, i: 1499] avg mini-batch loss: 0.765  
[epoch: 5, i: 2999] avg mini-batch loss: 0.732  
[epoch: 5, i: 4499] avg mini-batch loss: 0.749  
[epoch: 5, i: 5999] avg mini-batch loss: 0.757  
[epoch: 5, i: 7499] avg mini-batch loss: 0.772  
[epoch: 5, i: 8999] avg mini-batch loss: 0.781  
[epoch: 5, i: 10499] avg mini-batch loss: 0.753  
[epoch: 5, i: 11999] avg mini-batch loss: 0.764  
[epoch: 6, i: 1499] avg mini-batch loss: 0.678  
[epoch: 6, i: 2999] avg mini-batch loss: 0.728  
[epoch: 6, i: 4499] avg mini-batch loss: 0.733  
[epoch: 6, i: 5999] avg mini-batch loss: 0.738  
[epoch: 6, i: 7499] avg mini-batch loss: 0.754  
[epoch: 6, i: 8999] avg mini-batch loss: 0.731  
[epoch: 6, i: 10499] avg mini-batch loss: 0.735  
[epoch: 6, i: 11999] avg mini-batch loss: 0.760  
[epoch: 7, i: 1499] avg mini-batch loss: 0.676  
[epoch: 7, i: 2999] avg mini-batch loss: 0.685  
[epoch: 7, i: 4499] avg mini-batch loss: 0.693  
[epoch: 7, i: 5999] avg mini-batch loss: 0.714  
[epoch: 7, i: 7499] avg mini-batch loss: 0.708  
[epoch: 7, i: 8999] avg mini-batch loss: 0.725  
[epoch: 7, i: 10499] avg mini-batch loss: 0.730

[epoch: 7, i: 11999] avg mini-batch loss: 0.732  
[epoch: 8, i: 1499] avg mini-batch loss: 0.660  
[epoch: 8, i: 2999] avg mini-batch loss: 0.674  
[epoch: 8, i: 4499] avg mini-batch loss: 0.669  
[epoch: 8, i: 5999] avg mini-batch loss: 0.678  
[epoch: 8, i: 7499] avg mini-batch loss: 0.699  
[epoch: 8, i: 8999] avg mini-batch loss: 0.694  
[epoch: 8, i: 10499] avg mini-batch loss: 0.701  
[epoch: 8, i: 11999] avg mini-batch loss: 0.729  
[epoch: 9, i: 1499] avg mini-batch loss: 0.638  
[epoch: 9, i: 2999] avg mini-batch loss: 0.649  
[epoch: 9, i: 4499] avg mini-batch loss: 0.650  
[epoch: 9, i: 5999] avg mini-batch loss: 0.686  
[epoch: 9, i: 7499] avg mini-batch loss: 0.696  
[epoch: 9, i: 8999] avg mini-batch loss: 0.680  
[epoch: 9, i: 10499] avg mini-batch loss: 0.672  
[epoch: 9, i: 11999] avg mini-batch loss: 0.706  
[epoch: 10, i: 1499] avg mini-batch loss: 0.606  
[epoch: 10, i: 2999] avg mini-batch loss: 0.631  
[epoch: 10, i: 4499] avg mini-batch loss: 0.669  
[epoch: 10, i: 5999] avg mini-batch loss: 0.681  
[epoch: 10, i: 7499] avg mini-batch loss: 0.655  
[epoch: 10, i: 8999] avg mini-batch loss: 0.665  
[epoch: 10, i: 10499] avg mini-batch loss: 0.675  
[epoch: 10, i: 11999] avg mini-batch loss: 0.686  
[epoch: 11, i: 1499] avg mini-batch loss: 0.615  
[epoch: 11, i: 2999] avg mini-batch loss: 0.615  
[epoch: 11, i: 4499] avg mini-batch loss: 0.624  
[epoch: 11, i: 5999] avg mini-batch loss: 0.631  
[epoch: 11, i: 7499] avg mini-batch loss: 0.665  
[epoch: 11, i: 8999] avg mini-batch loss: 0.669  
[epoch: 11, i: 10499] avg mini-batch loss: 0.665  
[epoch: 11, i: 11999] avg mini-batch loss: 0.668  
[epoch: 12, i: 1499] avg mini-batch loss: 0.613  
[epoch: 12, i: 2999] avg mini-batch loss: 0.607  
[epoch: 12, i: 4499] avg mini-batch loss: 0.619  
[epoch: 12, i: 5999] avg mini-batch loss: 0.639  
[epoch: 12, i: 7499] avg mini-batch loss: 0.654  
[epoch: 12, i: 8999] avg mini-batch loss: 0.658  
[epoch: 12, i: 10499] avg mini-batch loss: 0.660  
[epoch: 12, i: 11999] avg mini-batch loss: 0.644  
[epoch: 13, i: 1499] avg mini-batch loss: 0.592  
[epoch: 13, i: 2999] avg mini-batch loss: 0.611  
[epoch: 13, i: 4499] avg mini-batch loss: 0.620  
[epoch: 13, i: 5999] avg mini-batch loss: 0.613  
[epoch: 13, i: 7499] avg mini-batch loss: 0.638  
[epoch: 13, i: 8999] avg mini-batch loss: 0.636  
[epoch: 13, i: 10499] avg mini-batch loss: 0.651

```
[epoch: 13, i: 11999] avg mini-batch loss: 0.680
[epoch: 14, i: 1499] avg mini-batch loss: 0.597
[epoch: 14, i: 2999] avg mini-batch loss: 0.616
[epoch: 14, i: 4499] avg mini-batch loss: 0.602
[epoch: 14, i: 5999] avg mini-batch loss: 0.610
[epoch: 14, i: 7499] avg mini-batch loss: 0.631
[epoch: 14, i: 8999] avg mini-batch loss: 0.606
[epoch: 14, i: 10499] avg mini-batch loss: 0.632
[epoch: 14, i: 11999] avg mini-batch loss: 0.664
Finished Training.
```

### Training Loss Curve

```
[8]: plt.plot(avg_losses)
plt.xlabel('mini-batch index / {}'.format(print_freq))
plt.ylabel('avg. mini-batch loss')
plt.show()
```



### Evaluate on Test Dataset

```
[9]: # Check several images.
dataiter = iter(testloader)
images, labels = next(dataiter)
imshow(torchvision.utils.make_grid(images))
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
outputs = net(images.to(device))
```

```
_, predicted = torch.max(outputs, 1)

print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                                for j in range(4)))
```



```
GroundTruth:   cat  ship  ship plane
Predicted:     cat  ship  ship plane
```

```
[10]: # Get test accuracy.
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (
    100 * correct / total))
```

Accuracy of the network on the 10000 test images: 70 %

```
[11]: # Get test accuracy for each class.
class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()
```



```
for i in range(4):
    label = labels[i]
    class_correct[label] += c[i].item()
    class_total[label] += 1

for i in range(10):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))
```

```
Accuracy of plane : 74 %
Accuracy of   car : 87 %
Accuracy of  bird : 65 %
Accuracy of   cat : 53 %
Accuracy of  deer : 60 %
Accuracy of   dog : 66 %
Accuracy of  frog : 77 %
Accuracy of horse : 69 %
Accuracy of  ship : 77 %
Accuracy of truck : 77 %
```

```
[ ]:
```